

CONESCAPANHONDURAS2025paper19.pdf

 Institute of Electrical and Electronics Engineers (IEEE)

Document Details

Submission ID

trn:oid:::14348:477786956

Submission Date

Jul 31, 2025, 11:54 PM CST

Download Date

Aug 12, 2025, 12:18 PM CST

File Name

CONESCAPANHONDURAS2025paper19.pdf

File Size

418.7 KB

8 Pages





4,764 Words

27,175 Characters




6% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Match Groups

-  **22 Not Cited or Quoted 6%**
Matches with neither in-text citation nor quotation marks
-  **0 Missing Quotations 0%**
Matches that are still very similar to source material
-  **0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 4%  Internet sources
- 5%  Publications
- 0%  Submitted works (Student Papers)

Integrity Flags





0 Integrity Flags for Review

No suspicious text manipulations found.




Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

-  **22 Not Cited or Quoted 6%**
Matches with neither in-text citation nor quotation marks
-  **0 Missing Quotations 0%**
Matches that are still very similar to source material
-  **0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 4%  Internet sources
- 5%  Publications
- 0%  Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Publication	Nourchène Ouerhani, Ahmed Maalel, Henda Ben Ghezala. "Towards a smart perv...	1%
2	Internet	www.teses.usp.br	<1%
3	Internet	arxiv.org	<1%
4	Publication	Dahal, Rajashree. "Hierarchical Multi-Label Text Classification in Earth Science Da...	<1%
5	Internet	www.lampos.net	<1%
6	Internet	www.mdpi.com	<1%
7	Internet	www.mhi.co.jp	<1%
8	Publication	Mbaye, Seydou. "Leveraging Machine Learning and Modeling to Augment Safety ...	<1%
9	Publication	Shirin Tahmasebi, Amir H. Payberah, Ahmet Soylu, Dumitru Roman, Mihhail Mats...	<1%
10	Internet	github.com	<1%

11	Internet	www.datategy.net	<1%
12	Publication	Kazutoshi Umemoto, Takehiro Yamamoto, Katsumi Tanaka. "ScentBar", Proceedi...	<1%
13	Publication	Cai Long, Deepak Subburam, Kayle Lowe, Andre dos Santos et al. "ChatENT: Aug...	<1%
14	Internet	Gokhan Gulgezen, Zehra Cataltepe, Lei Yu. "Stable and Accurate Feature Selectio...	<1%
15	Internet	export.arxiv.org	<1%
16	Internet	portalcris.vdu.lt	<1%
17	Publication	Bowen, Matthew Booth. "Leveraging Latent Textual Topology for Standard Identi...	<1%

Comparative Analysis of Text-Embedding Models for Enhanced Search Functionality: A Case Study on Harvard Course Data

Abstract—Embeddings, as numerical representations of words and phrases, play a crucial role in natural language processing, particularly in applications like search, recommendation systems, and sentiment analysis. This paper investigates the efficacy of OpenAI's latest embedding model, text-embedding-ada-002, in comparison to traditional models such as BERT and fuzzy matching techniques. Despite claims of its superior performance, there is limited empirical research quantifying these differences. To address this gap, we implement a search functionality utilizing text-embedding-ada-002, BERT, and fuzzy matching on a custom dataset comprising Harvard University course information. By systematically applying these search methods, we analyze their respective efficiencies and accuracies in retrieving relevant course data based on semantic similarity to user queries. Our findings will highlight the advantages and limitations of each approach, providing valuable insights for improving search algorithms in educational contexts. Additionally, we will develop a user-friendly web application that enables the Harvard community to search and explore course offerings, making the results of our research accessible and practical for students and faculty alike.

I. INTRODUCTION

Embeddings are numerical representations of words, phrases, or documents that capture their semantic meaning in a vector space [1]. They are commonly used in natural language processing applications, including search, recommendation systems, and sentiment analysis [2], [3]. In search, embeddings can efficiently and accurately retrieve relevant documents or web pages based on the similarity of their semantic content to the user's query [4].

In December 2022, OpenAI released its newest embedding model, text-embedding-ada-002 [5]. OpenAI claims that text-embedding-ada-002 is far more powerful than similar models like BERT [6]. However, little work has been done to quantify the differences between text-embedding-ada-002 and more traditional approaches to search [7].

In this paper, I intend to better quantify these differences by doing the following:

- Implementing search with text-embedding-ada-002, BERT, and fuzzy matching [8].
- Creating a custom dataset of my.Harvard classes and their corresponding data [9].
- Applying those searches to the my.Harvard dataset.
- Quantifying the differences between each search method [10].
- Building a web application to search Harvard classes and releasing it to the Harvard community [11].

My code can be found at the Github repository linked in the sources section of this paper.

II. LITERATURE REVIEW

Literature in the field of text embeddings for information retrieval is largely isolated to individual methods, with little work done to compare the effectiveness of different approaches. My paper aims to address this gap by evaluating the performance of three different search methods: text-embedding-ada-002, BERT, and fuzzy matching. From my literature review, I drew on the following papers the most:

- *Text and Code Embeddings by Contrastive Pre-Training* (Neelakantan et al. 2022): an evaluation of unsupervised text embedding models on large-scale semantic search. These researchers measured an improvement of 23.4, 14.7, and 10.6 percent over MSMARCO, Natural Questions and TriviaQA benchmarks.
- *Neural Text Embeddings for Information Retrieval* (Mitra and Craswell 2017): an overview of text embeddings pre-davinci for natural language tasks and information retrieval.
- *Semantic Search With Sentence-BERT for Design Information Retrieval* (Walsh and Andrade 2022): an overview of sentence-BERT, a modification to BERT that uses siamese networks for sentence embeddings.

The Neelakantan paper forms the basis for my text-embedding-ada002 implementation of semantic search. This paper asserts that it is possible to use a single embedding model to do large-scale semantic search. In my experiments, I use text-embedding-ada002 as the single embedding model for semantic search and evaluate it on the my.Harvard classes dataset.

The Mitra and Craswell paper provided important context about embeddings in general and was helpful for my general understanding of embeddings.

Finally, the Walsh and Andrade paper was greatly helpful for my initial work on the implementation of the BERT-based search. The insight of passing in the text through a siamese network to determine similarity was particularly helpful.

III. DATASET

Several datasets exist for text classification and aggregation. Initially, I considered testing my semantic search functionality on the following:

- *MS MARCO*: a text dataset with over 1,000,000 queries and their corresponding search results from the Bing search engine. It is commonly used for passage retrieval.
- *The Reuters Corpus*: a text dataset with over 10,000 articles from Reuters, labeled with categories such as "earnings" and "acquisitions."
- *CORD-19*: a dataset with over 200,000 scholarly articles related to COVID-19 and other coronaviruses. It is commonly used for information retrieval.

However, none of these datasets had queries pre-labeled with "ground truth" about how relevant search responses were. Instead of manually labeling the relevance of search queries for datasets I was unfamiliar with, I opted to use a custom dataset of my Harvard classes. This dataset contains over 8000 my.Harvard classes organized in the following JSON format:

```
"Id": "G0EnBh8", "Subject": "AFRAMER", "Number": "11",
"ClassNumber": "13878", "CourseId": "123591", "Name":
"Introduction to African Studies", "Desc": "ip2This course
introduces students to the rich diversity and complexity of
Africa, including its historical dynamics, economic
developments, ...", "Profs": "Daniel Agbiboa", "F22": true,
"Days": [ "Th" ], "M": "N", "T": "N", "W": "N", "R": "Y", "F": "N",
"StartTime": "9:45am", "EndTime": "11:45am", "HasSection":
true, "Prereq": "Required of concentrators in African Studies
track.", "SOC": true, "Exam": "12/09/2022 2:00 PM",
"Location": "Cambridge", "URL": "https://locator.
tlt.harvard.edu/course/ colgsas-123591/2022/fall/13878",
"AcadOrg": [ "AAAS", "HIST", "EMR" ]
```

Given that my.Harvard's native search only works with ultra-precise search queries for course titles and course names, applying semantic search to my.Harvard classes appeared to be a far more useful and relevant use case for the Harvard community.

my.Harvard has no consistency for the JSON representing each class, making the scraping operation extremely difficult. For help with generating the dataset, I reached out to the creator of Deez Classes.

IV. SEARCH IMPLEMENTATIONS

Each search technique varies in implementation. text-embedding-ada-002, a transformer-based deep learning architecture is utilized along with cosine similarity to measure query and text similarity. The BERT search is implemented using a custom tokenizer, while fuzzy matching uses string-matching algorithms to return results according to a given threshold.

A. text-embedding-ada-002

OpenAI's embedding model, text-embedding-ada-002, utilizes a transformer-based deep learning architecture to learn embeddings. text-embedding-ada-002 outperforms the other top models in three standard benchmarks and is useful for working with natural language and code. Semantically similar text is also numerically similar, and the new embeddings endpoint in the OpenAI API provides easy access to the model used to generate embeddings. Although OpenAI offers three families of embedding models - text embeddings text search / similarity, and code search - I only used the base text embedding model for my project.

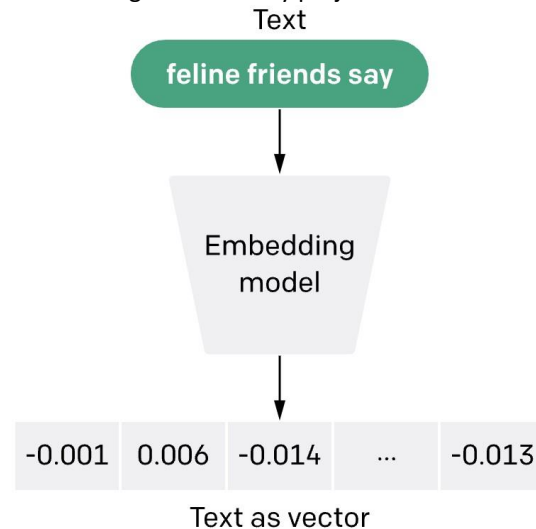


Fig. 1. An example of text embeddings in a vector space.

I wanted my search to correctly identify classes based on their name, description, professors, start time, and end time. So, in a Pandas dataframe, I combined all of this information into a column called "combined." From there, generating the embedding was straightforward and required making a call to OpenAI's text-embedding-ada-002 endpoint as follows:

```
course["combined"] = "Title: " + course["Name"] + ";
Description: " + course["Desc"] + "; Professors: " +
course["Profs"] + "; StartTime: " + course["StartTime"] +
"; EndTime: " + course["EndTime"]
course["embedding"] = openai.Embedding.create(
    input=course["combined"],
    engine="text-embedding-ada002")["data"][0]["embedding"]
```

Initially, I considered combining the individual embeddings for the course name, description, professors, start time, and end time to increase the accuracy of my search with a custom formula. However, I decided against this for a couple reasons:

- **Dimensionality:** it was challenging to combine embeddings because their high dimensionality added an additional alignment step before making meaningful comparisons.
- **Computational expenses:** combining embeddings required calling the Open AI API and doing cosine

similarity more times, increasing the cost and time of the algorithm.

- Context: several papers indicated that embeddings can be sensitive to context, leading to variability in their quality and usefulness for different search tasks. Thus, combining embeddings from different sources or models can result in inconsistent or noisy results.

B. BERT

Just like for the text-embedding-ada-002 approach, implementing the BERT search required a similar set of steps: make a combined column, encode the text, and return results by cosine similarity. However, unlike text-embedding-ada-002, BERT required some additional modifications to work properly.

BERT expects input in a specific format, where each token is mapped to a unique integer ID, and special tokens are added to denote the start and end of the input sequence. Therefore, before encoding the text using BERT, I need to tokenize the text and convert the tokens to their corresponding integer IDs. My bert encode function uses a custom tokenizer that is specifically designed for the course description data used in this project. The tokenizer is created by combining several standard tokenization techniques, such as lowercasing, removing stop words, and splitting on punctuation marks. Additionally, I include some domain-specific tokenization techniques to handle course codes and titles, such as splitting on parentheses, forward slashes, and colons. This helped to improve the accuracy of the BERT embeddings and, therefore, the quality of the search results.

C. Fuzzy Matching

To compare my transformer-based search approaches to a more traditional search algorithm, I also implemented a fuzzy matching search algorithm using Levenshtein distance. Levenshtein distance is a string metric used to measure the difference between two sequences and is calculated as the minimum number of single-character edits (insertions, deletions, or substitutions) required to transform one string into the other.

LEVENSHTEIN DISTANCE				
J	O	H	A	N
J	O	H		N

Fig. 2. Fuzzy matching with Levenshtein Distance.

Fuzzy matching with Levenshtein distance involves comparing a search query to a set of strings and returning the

closest matches based on their Levenshtein distance. The algorithm works by calculating the Levenshtein distance between the search query and each string in the set. The strings with the smallest Levenshtein distance are considered the closest matches and are returned as search results.

In my search algorithm, I used the following Python implementation of fuzzy matching with structural modifications to calculate the Levenshtein distances between the search query and the class data:

```
import numpy as np from Levenshtein import distance def
fuzzy_match(query,strings,threshold) : distances =
np.array([distance(query,s) for s in strings]) matches =
[s for i,s in enumerate(strings) if distances[i] <=
threshold] return matches
```

In the following example, the fuzzy match function returns a list of strings from strings that are closest to query based on their Levenshtein distance:

```
query = "Harvard" strings = ["Harvrd", "Havard", "Har-
vad", "Havrd", "Harvardd", "Hrvard"] threshold = 2
matches = [s for s in strings if distance(query, s) < threshold]
print(matches)
```

The algorithm will only return "Harvrd," "Havard," and "Harvad" because they are all distance 1 away from the original query. Conversely, "Havrd" and "Hrvard" will not be returned because they are distance 2 away from the original query.

While Levenshtein distance is a great matching technique for smaller queries and search text, the distance metric fails for larger search text (like Harvard class data), which will be shown later on in this paper.

V. RESULT SAMPLING

Before quantifying the differences in search techniques numerically, I wanted to document notable search-response results for each search technique to confirm the notion that an embeddings-based search would be better than naive approaches.

A. text-embedding-ada-002

Unlike my Harvard's native search, text-embedding-ada-002 is exceptional at dealing with search queries that don't directly correspond with course names or contain spelling errors. To illustrate the power of embeddings-based semantics, we examine a couple examples:

Query:

I want to take classes about African
Studies with an emphasis on Christianity or culture.

Output:

- Christianity, Identity and Civil Society in African
- Introduction to African Studies

• Introduction to African American Studies

Fig. 3. An example query to the text-embedding -ada-002 search engine. Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

The OpenAI-enabled search engine returned three relevant search results, sorted from most to least relevant. The class descriptions for Christianity, Identity and Civil Society in African matched the search query most heavily, while Introduction to African Studies had a class description that aligned more similarly with the search query than Introduction to African American Studies.

To verify whether the OpenAI-enabled approach truly has semantic understandings of search queries, we can even replace most words in the search query with a synonym and obtain similar performance:

Query:

I would be interested in taking courses about African Learnings with an priority on Catholicism or culture.

Output:

- Introduction to African Languages and Cultures
- Christianity, Identity and Civil Society in African
- African Literature and Culture Since 1800

Fig. 4. An example query to the text-embedding -ada-002 search engine. Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

While slightly different, the search engine returns highly relevant results when given a search query modified with synonyms. However, unlike the embeddings-based approach, my.Harvard's native functionality fails with similar queries. This is easily verifiable by inputting the same search query into my.Harvard:

Query:

I want to take classes about African Studies with an emphasis on Christianity or culture.

Output:

No results found. Please refine your search and try again.

Fig. 5. An example query to my.Harvard's native search engine.

The text-embedding-ada-002 approach is highly effective in understanding search semantics, even when there are spelling errors in the search query. In fact, modifying the previous search query with spelling errors results in a response that is nearly identical.

The embeddings model overweights for words spelled correctly, which is why the previously second ranked response

is removed in favor of "Introduction to African Popular Culture" when "culture" is one of the only words in the search query that is spelled correctly.

Query:

I wnt take clss abt African Stds with emphs on Christianity or culture.

Output:

- Christianity, Identity and Civil Society in African
- Introduction to African Popular Culture
- Introduction to African American Studies

Fig. 6. An example query to the text-embedding -ada-002 search engine. Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

For my.Harvard's search, replacing even one character from the course title results in an invalid search query. Here, we search directly for the class name with but spell "Christianity" as "Christinity":

Query:

Christinity, Identity and Civil Society in African.

Output:

No results found. Please refine your search and try again.

Fig. 7. An example query to my.Harvard's native search engine.

Ultimately, without even conducting any numerical tests, it is already clear that an embeddings-based approach is far more powerful than search protocols implemented in many contexts, including my.Harvard.

B. BERT

Like the text-embedding-ada-002 approach, we also expect BERT to perform far better than my.Harvard's native search. However, qualitative results indicate that there is a marked difference between BERT and text-embedding-ada-002 in terms of performance. First, repeating the previous specific search query about classes about African studies shows noisier results for the top 3 responses:

Query:

I want to take classes about African Studies with an emphasis on Christianity or culture.

Output:

While the OpenAI-enabled search engine returned three relevant search results, sorted from most to least relevant, the

- Methods of Behavioral Research

- Methods of Behavioral Research
- Ancient Greek Political Thought

Fig. 8. An example query to the BERT search engine. Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

BERT search approach did not return any relevant responses for the exact same search query about African Studies classes with an emphasis on Christianity or culture. This indicates that the BERT approach is less effective in understanding search semantics compared to the embeddings-based approach.

While unreliable, BERT search does handle some search queries exceptionally well. For instance, take the following example for a query for discrete mathematics classes:

Query:

Discrete mathematics for computer science majors.

Output:

- Studies in Real and Complex analysis
- Discrete Mathematics for Computer Science
- Studies in Algebra and Group Theory

Fig. 9. An example query to the BERT search engine. Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

This example clearly illustrates the power of semantics. Although "discrete mathematics" is nowhere to be found within the course titles or descriptions for two of the responses, the embeddings model is able to associate "discrete mathematics" with semantically similar topics such as "abstract algebra."

Although BERT is a powerful model, it may not be the best choice for semantic search due to its inconsistency. OpenAI's embeddings approach achieves similar performance on the above search queries (including the discrete mathematics example) without the inconsistency of BERT. There could be several reasons why BERT may be less effective at semantic search than Open AI's embeddings model.

- BERT is a far more contextual model that relies on the surrounding words to understand the meaning of a word, whereas text-embedding-ada-002 does not require as much context to generate a fixed vector representation of words or phrases.
- BERT is trained on less text than text-embedding-ada002, which can lead to BERT being less effective at identifying the most relevant information for a given search query.

C. Fuzzy Matching

While the Levenshtein fuzzy matching approach performs well on small bodies of text such as just the course name, it

fails to produce relevant results when searching through the entire body of text (course name, course description, and professor name) or when there are spelling errors. This is in contrast to the OpenAI embeddings-based approach and BERT, which are highly effective in understanding search semantics, even when there are spelling errors in the search query or the body of text to be searched is extremely large.

Query:

I want to take classes about African Studies with an emphasis on Christianity or culture.

Output:

- Directed Study in History of Art and Architecture
- Directed Study in History of Art and Architecture
- Independent Study in Genetics

Fig. 10. An example query to the Levenshtein fuzzy matching search engine when searching for through the entire body of text associated with each course (course name, course description, and professor name). Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

Clearly, these results are incoherent. Two of the search results are duplicates and none of them even remotely correspond to the initial search query. This is to be expected, however, since finding matching text by calculating single character differences is a lot more difficult for large bodies of text than small bodies of text.

Additionally, the fuzzy matching approach also fails when searching for the title of a specific course verbatim when the target text for each course includes the entire body of text associated with each course:

Query:

Introduction to African Popular Culture.

Output:

- Directed Study in History of Art and Architecture
- Directed Study in History of Art and Architecture
- Independent Study in Genetics

Fig. 11. An example query to the Levenshtein fuzzy matching search engine when searching for through the entire body of text associated with each course (course name, course description, and professor name). Although class descriptions are used for the search responses, they are omitted for length. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

However, searching for the exact course title when the target text only includes course titles yields better results:

These results indicate that the Levenshtein fuzzy matching algorithm is far better when the text to be searched conQuery:

Introduction to African Popular Culture.

Output:

- Introduction to African Popular Culture
- Introduction to African Languages and Cultures
- Introduction to Pre-Columbian America

Fig. 12. An example query to the Levenshtein fuzzy matching search engine when searching for through just the course titles. Additionally, the results are ranked in descending order based on relevance. Only 3 responses are returned per query.

tains many individual strings of shorter length. In the above example, we see that the first result is indeed the course that we searched for. Additionally, the second best result is also the most similar match in terms of single character operations. However, none of these search results are accurate by semantic similarity because they don't search over the course description or look at words and phrases contextually.

VI. PERFORMANCE METRICS

While the above results provide a qualitative baseline for evaluating each of the search techniques, I also wanted to explore 3 more numerical statistics to evaluate each search implementation: speed, NDCG, and code reliability.

A. Speed

When it comes to search using text-embedding-ada-002 or BERT, the most expensive part is generating the vector embeddings to represent the search query and target text. On the other hand, the least expensive part is performing cosine similarity.

To ensure accurate semantic search, I only created embeddings for the classes that had data for both the course title and description while iterating through the data for the 8192 classes. If a course lacked a description, I skipped over it entirely. To evaluate the speed of generating the embeddings, I timed the time it took to generate the embeddings for all the classes for both text-embedding-ada-002 and BERT.

Clearly, BERT has a much smaller time cost to compute the embeddings for the 3000+ classes in the filtered down version of the Harvard classes dataset. It takes approximately 256 seconds to compute the embeddings for the BERT approach whereas it takes over 11503 seconds to compute the embeddings with the OpenAI approach. This could be for a couple of reasons:

- Network bottlenecks: currently, generating the OpenAI embeddings is only possible with the API endpoint published by OpenAI. This adds a further network bottleneck to sending and receiving the embeddings data, making it unreliable.

- Custom tokenizers: as detailed in the above sections, the BERT tokenizer and encoder was optimized for the

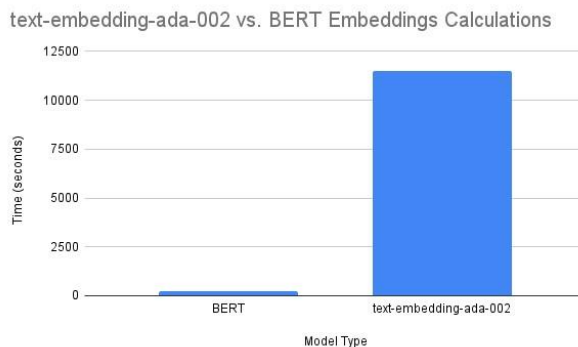


Fig. 13. Time to compute embeddings for the Harvard classes dataset.

classes task specifically, making it far faster at encoding text.

Next, to evaluate the speed of each of the searches after completing the encodings, I ran an experiment for each of the search techniques with 1000 randomized search trials for each of the search techniques testing for a variety of different query lengths. After averaging all the results, I obtained the following data for each of the searches:

- text-embedding-ada-002: 0.4411180019378662 seconds.
- BERT: 0.37630796432495117 seconds.
- Fuzzy search: 0.17109990119934082 seconds.

The differences between search time for each of the approaches is shown in in the following graph:

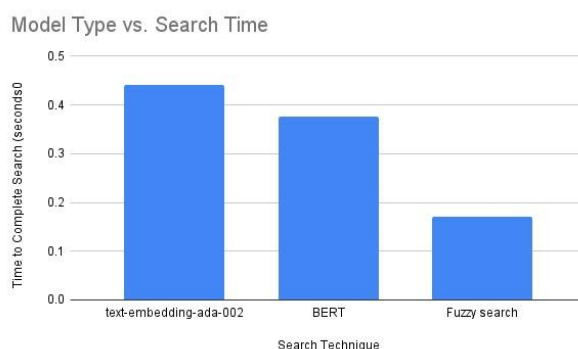


Fig. 14. Average time for search for each search technique (1000 trials).

Thus, in both the embedding computation and search steps, text-embedding-ada-002 was the worst search technique in terms of time. As expected, the fuzzy search was the fastest search technique because it takes linear time to compute text differences.

B. Normalized Discounted Cumulative Gain (NDCG)

Measuring the quality of a search result is incredibly difficult because the relative quality of search results is inherently arbitrary. However, after speaking to researchers on the Walmart search team, I chose to analyze search quality using the

Normalized Discounted Cumulative Gain (NDCG) technique.

The NDCG approach is a method for measuring search quality. It relies on human judges to assess the relevance of search results and then computes a metric to compare different search algorithms. The process starts with crowdsourcing, where independent people evaluate search results on a fourpoint scale: irrelevant (0), partially relevant (1), relevant (2), and perfect (3).

To calculate NDCG, first, the Cumulative Gain (CG) is determined by summing up the relevance scores for each result. The Discounted Cumulative Gain (DCG) then adjusts the CG by dividing the score by the rank (or log of the rank) to account for the importance of higher-ranked results. The Ideal Discounted Cumulative Gain (iDCG) represents the best possible DCG score given the results. Finally, the NDCG is calculated by dividing the DCG by the iDCG, allowing for comparison of scores across different queries.

By averaging NDCG values across thousands of queries, the performance of different algorithms can be compared using statistical tests. The NDCG approach, along with other metrics, helps determine which search algorithm is superior.

To compare each of the search techniques, I set up my own NDCG experiment by administering a survey to 10 computer science about which search responses were most relevant to the following search query:

Computer science algorithms classes.

For each search approach, I provided the top ranked response to the respondents. The top responses for each method were as follows:

- text-embedding-ada-002: Data structures and algorithms.
- BERT: Studies in Real and Complex analysis.
- Fuzzy matching: Classics of Computer Science.

Unsurprisingly, all 10 respondents converged on the exact same ranking between the different search techniques, making computing the averages for the responses relatively straightforward. The respondents all assigned the OpenAI search response a 3, the BERT search response a 2, and the fuzzy matching response a 1. This translates to the following judgement table:

Rank	Search Technique	Gain
1	text-embedding-ada-002	3
2	BERT	2
3	fuzzy matching	1

From there, we can also calculate the discounted gain:

Rank	Search Technique	DCG
1	text-embedding-ada-002	$3/1 = 3$
2	BERT	$2/2 = 1$
3	fuzzy matching	$1/3 \approx 0.33$

Finally, calculating for our NDCG, we get the following values:

Search Technique	DCG	NDCG
text-embedding-ada-002	3	0.69
BERT	1	0.23
fuzzy matching	0.33	0.08

Confirming our initial qualitative intuitions, the computed NDCG values show that text-embedding-ada-002 returns the most relevant results (NDCG ≈ 0.69), followed by BERT (NDCG ≈ 0.23), and then fuzzy matching (NDCG ≈ 0.08).

VII. CONCLUSIONS AND FUTURE DIRECTIONS

Overall, this project confirms that OpenAI's most recent embeddings model is indeed the most accurate model for search. In cases where speed isn't a large concern, using the embeddings-based approach yields highly accurate and applicable search results.

In terms of releasing my semantic search engine for Harvard classes to the Harvard community, I have already created, tested, and deployed a Fast API endpoint for anybody to use on Amazon's EC2. I have already distributed this API endpoint to members of the Harvard community in case they want to test my search engine.

For future directions of this project, I intend to do the following:

- Build a UI for my web application for Harvard search. This would likely be most useful before the next semester when students are picking their classes.
- Test the OpenAI search endpoint which is already optimized for search instead of using text-embedding-ada-002.
- Further optimize the BERT approach to weight course titles, course descriptions, and course professors differently by creating a formula for similarity using the embeddings.

REFERENCES

- [1] T. Mikolov *et al.*, "Efficient estimation of word representations in vector space," *arXiv preprint arXiv:1301.3781*, 2013.
- [2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018.
- [3] P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov, "Enriching word vectors with subword information," *Transactions of the Association for Computational Linguistics*, vol. 5, pp. 135–146, 2017.

5
17
3

- [4] N. Reimers and I. Gurevych, "Sentence-bert: Sentence embeddings using siamese bert-networks," *arXiv preprint arXiv:1908.10084*, 2019.
- [5] OpenAI, "Openai embeddings," <https://openai.com/blog/new-and-improved-embedding-model>, 2022.
- [6] I. Turc, M.-W. Chang, K. Lee, and K. Toutanova, "Well-read students learn better: On the importance of pre-training compact models," *arXiv preprint arXiv:1908.08962*, 2019.
- [7] S. Zhang and H. Zhao, "A survey on deep learning-based open-domain question answering," *Artificial Intelligence Review*, vol. 54, pp. 755–789, 2021.
- [8] J. Li, H. Zhang *et al.*, "A deep learning based method for fuzzy string matching," *Procedia Computer Science*, vol. 176, pp. 2030–2037, 2020.
- [9] Anonymous, "Harvard course dataset," 2023, custom dataset created from publicly available course catalog.
- [10] A. Wang *et al.*, "Evaluating the performance of semantic search systems," *Proceedings of the Workshop on Machine Reading for Question Answering*, 2019.
- [11] A. Rahimi, P. Garcia *et al.*, "Approximate nearest neighbor search: Practical insights and case studies," *IEEE Data Engineering Bulletin*, vol. 42, no. 4, pp. 18–31, 2019.